

Convolution and CNNs

Terms you may run into

CNN (Convolutional Neural Network):


- A deep learning model type designed to process **grid-like** data, such as images. CNNs use convolutional layers to automatically detect important features in images, like edges, shapes, and textures, making them effective for image classification and object detection.

Pooling:

- A down-sampling operation in CNNs that reduces the spatial dimensions (width and height) of feature maps, making the model more efficient and resilient to small image distortions. Common types are max pooling (taking the maximum value in a region) and average pooling (taking the average).

LeNet:

- One of the earliest CNN architectures, developed by Yann LeCun for handwritten digit recognition (MNIST dataset). It introduced convolutional and pooling layers and laid the groundwork for future CNNs.



Edge Detection:

- A process in image processing that identifies the boundaries or edges within an image, which are areas with significant changes in brightness or color. Edge detection highlights structures like object outlines, making it useful for detecting shapes and patterns. In CNNs, certain filters (kernels) are applied to input images to perform edge detection.

Kernel:

- A small matrix or "window" of numbers (like a 3×3 matrix) used in convolution operations in CNNs. The kernel slides over the image, performing element-wise multiplications and summing them up to detect specific features (like edges or textures) at each position. Different kernels are used to detect various features.

Filter:

- In CNNs, filters (or **kernels**) are used interchangeably. A **filter** is essentially a kernel applied to an input image to detect specific features, such as edges, textures, or shapes. A CNN might use multiple filters in each layer to detect a wide range of features at different levels of abstraction (edges, corners, etc.).

Feature Map:

- A **feature map** (also called an **activation map**) is the output of a convolutional layer after a filter (or kernel) has been applied to the input. It highlights certain **features** in the input image, for example, edges, corners, textures, or patterns.
- A **filter** is a trained detector that looks for a specific type of feature.
- A **feature map** is the result of sliding that filter over the entire input — showing where that feature is present in the input.
- The number of **feature maps** generated by a convolutional layer is **equal to the number of filters** used in that layer.

Example

If a layer has **32 filters**, each looking for a different feature...

Then the output of that layer will be **32 feature maps** — one for each filter.

What does each feature map represent?

Each feature map corresponds to:

- A **specific kind of pattern** (like vertical edges, diagonal lines, textures, color transitions).
- The **spatial layout** of where that pattern shows up in the input.

As you go deeper in the network:

- **Early layers** detect **simple features** (e.g., edges, gradients).

- **Later layers** detect **complex features** (e.g., faces, objects, textures).

Dimensions of a Feature Map

- If the input is 64×64 and the convolution uses:
 - kernel size 3×3
 - stride = 1
 - padding = "same"

Then the output feature map will still be roughly 64×64 in spatial size (but smaller without padding).

- If you have 32 filters, the total output will be a tensor of size **$64 \times 64 \times 32$** — i.e., 32 **feature maps**, one per filter.

Example

```
nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3)
```

- You are applying **64 filters**, each producing **1 feature map**.
- The output will be **64 feature maps** (also referred to as **channels** in the output tensor).

Data Augmentation:

- Techniques, such as random cropping and flipping, to artificially increase the size of the dataset and make the model more resilient to variations.

Architectures

1. Image Classification Networks

- **VGGNet (2014)**: A deep network with sequential layers, known for its simplicity and use of 3×3 filters. It's straightforward but computationally intensive.

- **GoogLeNet / Inception (2014/2015)**: Introduced the Inception module, which uses multiple filter sizes in parallel, reducing computational costs by focusing on efficient architecture.
 - **ResNet (2015)**: Uses residual connections (skip connections) to solve the vanishing gradient problem, allowing very deep networks by adding identity mappings.
 - **Xception (2016)**: An extension of Inception, with depthwise separable convolutions that help to improve computational efficiency while maintaining accuracy.
 - **MobileNet (2017)**: Lightweight network focused on mobile and embedded vision applications, using depthwise separable convolutions to reduce parameters.
 - **EfficientNet (2019)**: Uses a compound scaling method that optimally balances model depth, width, and resolution, achieving high accuracy with fewer parameters.
-

2. Object Detection Networks

- **YOLO (You Only Look Once) (2016)**: Real-time object detection network that splits the image into a grid, assigning bounding boxes and class probabilities for each grid cell, achieving fast inference.
 - **SSD (Single Shot MultiBox Detector) (2016)**: Uses multiple bounding boxes per feature map location, enabling multi-scale detection and handling objects of varying sizes.
 - **Faster R-CNN (2015)**: A two-stage network that generates region proposals (regions likely to contain objects) and then classifies and refines these regions, leading to high accuracy but slower inference.
-

3. Instance Segmentation Networks

- **Mask R-CNN (2017)**: Extends Faster R-CNN by adding a branch for predicting segmentation masks on each region of interest, making it suitable for pixel-level segmentation of individual instances.
-

4. Semantic Segmentation Networks

- **U-Net (2015):** Originally designed for biomedical image segmentation, it uses an encoder-decoder architecture with skip connections to capture fine details in segmentation.
- **SegNet (2015):** An encoder-decoder architecture that stores max-pooling indices to improve memory efficiency and accuracy in semantic segmentation tasks.
- **DeepLab (2016):** Uses atrous convolutions (dilated convolutions) and an improved backbone to capture larger context, improving segmentation in complex scenes.
- **FCN (Fully Convolutional Network) (2015):** A network that replaces dense layers with convolutions, allowing it to output pixel-level segmentation by upsampling lower-resolution outputs.

Questions and answers

Is there a minimum basic structure for something to be considered and function as a cnn for classification?

Yes, there is.

- **Convolutional Layer:** At least one Conv2D layer to extract features
 - Learns spatial patterns in the input
 - Defined by kernel size, number of filters, stride, and padding
- **Activation Function:** Typically ReLU
 - Introduces non-linearity
 - Helps learn complex patterns
- **Pooling Layer:** Usually MaxPool2D
 - Reduces spatial dimensions
 - Makes the model more robust to position variations
- **Fully Connected Layer:** For final classification
 - Converts features to class probabilities

- Output size matches number of classes

Are convolution networks just using deep learning for finding filters (kernels, feature detectors)?

Yes.

Longer answer

- **Learning Filters:** Traditional image processing often uses hand-crafted filters (like edge detectors) designed by experts. In contrast, CNNs learn these filters from data, which allows the network to capture relevant features directly from images. This process is part of the training, where backpropagation adjusts the filter values to minimize the loss function, so they become optimized for the specific task, such as classifying objects or detecting edges.
- **Hierarchical Feature Extraction:** In a CNN, earlier layers often learn simpler patterns like edges or textures. As you go deeper, layers learn increasingly complex and abstract features, combining the simpler patterns learned in previous layers.
- **Automatic Filter Discovery:** During training, CNNs explore thousands of potential filter combinations and orientations, allowing them to find the most effective filters for the task, even those that might be unintuitive or difficult for a human to design.

Why do we need pooling?

Pooling reduces the spatial dimensions of feature maps, keeping essential features while making the network more efficient.

Why do we need padding?

Padding adds extra pixels around the input to keep output dimensions consistent or control the spatial size after convolution.

Does convolution change dimensions?

Convolution reduces the spatial size of the output based on the filter size and stride, so padding helps control this reduction by adding pixels around the edges.

But isn't convolution an operation on two functions that produces a third function?

- **Mathematical Definition:** In traditional mathematics and signal processing, convolution is an operation on two functions that produces a third function. This operation essentially blends one function with another, showing how one modifies the other over a range.
- **CNN Adaptation:** In convolutional neural networks, this idea is **adapted** to data like images. Here, convolution involves sliding a filter (a small matrix) over an image to perform dot products at each position. While it's not the exact same as the full mathematical definition, it retains the essence of combining input data with a filter to produce a new "function" or feature map. This feature map captures important information for the network to learn from.

Where can i find image datasets?

I would start with Roboflow

- is well integrated with YOLO
- <https://universe.roboflow.com/>

Then hugging face

<https://huggingface.co/datasets>

A bit unorganized, but has image datasets.

<https://images.cv/>